



## SODA 24

A brief introduction to some interesting papers

Zhiwei Ying  
2024-08-07



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY

# 目录



- 1 Other topics
- 2 Online Algorithms

1

## Other topics

Shortest Disjoint Path on a Grid

2

## Online Algorithms



## Other topics

Shortest Disjoint Path on a Grid



## Online Algorithms

# Shortest Disjoint Path on a Grid




# Settings



Given a graph  $G$  on a grid, and  $2k$  points  $s_1, t_1, \dots, s_k, t_k$ . The goal is to find  $k$  vertex-disjoint paths that joins  $s_i, t_i$  for  $i \in [k]$  with the total length of the solution to be minimized.



Given a graph  $G$  on a grid, and  $2k$  points  $s_1, t_1, \dots, s_k, t_k$ . The goal is to find  $k$  vertex-disjoint paths that joins  $s_i, t_i$  for  $i \in [k]$  with the total length of the solution to be minimized.



# Other Settings



- 1 The problem that given pairs of terminals with integer coordinates on a infinite grid, the player only need to decide whether vertex disjoint paths connecting these terminals exist is proved to be NP-Complete.



# Other Settings



- ① The problem that given pairs of terminals with integer coordinates on a infinite grid, the player only need to decide whether vertex disjoint paths connecting these terminals exist is proved to be NP-Complete.
- ② edge-disjoint
- ③ on directed graph or general planar graph
- ④ with length bounds on paths
- ⑤ Given a set of terminals, return the largest value  $k$  that there exist  $k$  vertex-disjoint paths joining a pair of terminals.




Given a set  $\Gamma = \{(s_1, t_1), \dots, (s_k, t_k)\}$  of  $k$  pairs of terminals placed arbitrarily on a grid, we can find the shortest  $k$  vertex-disjoint paths between  $\Gamma$  or attest that there is no feasible solution in time  $k^{2^{O(k)}} O([\Gamma])$ , where  $[\Gamma]$  is the number of bits required to encode  $\Gamma$  in binary.



**Irrelevant vertices:**

- ❶ outside a bounding box of the terminals, whose border is at distance  $\Theta(2^k)$  from the extreme terminals.
- ❷ inside a bounding box of the terminals, whose border is at distance  $\Theta(k \cdot 2^k)$  from the horizontal and vertical lines containing the terminals.



# Main Technique



## Irrelevant vertices:

- ① outside a bounding box of the terminals, whose border is at distance  $\Theta(2^k)$  from the extreme terminals.
- ② inside a bounding box of the terminals, whose border is at distance  $\Theta(k \cdot 2^k)$  from the horizontal and vertical lines containing the terminals.

After excluding these irrelevant vertices, the number of solutions to enumerate is reduced to  $\Theta(k \cdot 2^k \cdot L)$ , where  $L$  is the distance between 2 terminals.

## Structural lemma:


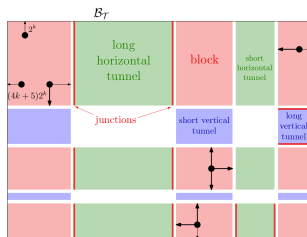
If there exists a routing of a set of terminals  $\Gamma$ , then there exists an optimal routing of  $\Gamma$  that does not use any irrelevant vertex.

# Main Technique



## Key lemma:

For those long horizontal tunnels that whose width is greater than its height, denote  $J_l$  and  $J_r$  as their junctions, and  $\hat{J}_l = \{v_1, \dots, v_t\} \subseteq J_l$  ordered from top to bottom,  $\hat{J}_r = \{v'_1, \dots, v'_t\} \subseteq J_r$  ordered from top to bottom. Then there exists a family  $(Q_j)_{j=1}^t$  of pairwise disjoint paths.



# Algorithm



Map  $V'$  that consists of all vertices in a block, in a short tunnel, or in a junction of a long tunnel to  $\{0, 1, \dots, k\}$ . Enumerate all mappings and then get the result.

## Running time:

The size of  $V'$  is  $2^{O(k)}$ . The number of all possible mappings is  $(k+1)^{2^{O(k)}}$ . All vertices in  $V'$  are contained in a bounding box that is at a distance at most  $2^k$  from the coordinates of each vertex and can be encoded using  $O(2(b + \log(2^k))) = O([\Gamma])$  bits, where  $b$  is the maximum number of bits of a coordinate of a point in  $\Gamma$ . And we can decide in time  $\text{poly}(2^{O(k)})$  whether this configuration is valid or not.

Thus the overall running time is  $k^{2^{O(k)}} O([\Gamma])$ .

# Open Problems



What will happen if more points are given? Given  $k$  sets of points  $S_1, \dots, S_k$  on a grid with  $\forall i \in [k], |S_i| \geq 2$ , how to find  $k$  vertex-disjoint paths that connects each points in a set with the total length of the solution to be minimized?



Other topics



## Online Algorithms

Online Robust Mean Estimation

Dynamic algorithms for k-center on graphs

Bin Packing under Random-Order: Breaking the barrier of  $\frac{3}{2}$





Other topics



**Online Algorithms**

Online Robust Mean Estimation

Dynamic algorithms for k-center on graphs

Bin Packing under Random-Order: Breaking the barrier of  $\frac{3}{2}$

# Online Robust Mean Estimation



# Background



Consider a general offline setting first. Given  $n$  samples from a distribution  $\mathcal{X}$  on  $\mathbb{R}^m$  with unknown mean  $\mu^*$ . The algorithm should return an estimate  $\hat{\mu}$  of  $\mu^*$ , and make  $\|\hat{\mu} - \mu^*\|$  as small as possible.

However, as is known to all, some samples are generated by malicious users, especially competitors, and thus a fraction of data might be misleading or meaningless. A formal definition of this process is given below.

## Definition 1.1.1 ( $\varepsilon$ -corrupted)

Given a parameter  $0 \leq \varepsilon \leq \frac{1}{2}$  and a set  $\mathcal{C}$  of  $n$  samples, an  $\varepsilon$ -corrupted version of  $\mathcal{C}$  is generated after an adversary removes up to  $\varepsilon n$  samples from  $\mathcal{C}$  and replace them by arbitrary points.

# Settings



Given  $M, T \in \mathbb{Z}^+$ , s.t.  $M$  is an interger multiple of  $T$ ,  $0 \leq \varepsilon < \frac{1}{2}$ , let  $\mathcal{X} = \{x^{(1)}, \dots, x^{(n)}\}$  be an  $\varepsilon$ -corrupted version of a clean set of *i.i.d.* samples from a distribution  $X$  on  $\mathbb{R}^M$  with unknown mean  $\mu^*$ . The  $M$  coordinates of each data point are divided into  $T$  batches, each of size  $d \triangleq \frac{M}{T}$ , i.e.  $x^{(i)}$  is the concatenation of  $x_1^{(i)}, \dots, x_T^{(i)}$ , where  $x_t^{(i)} \in \mathbb{R}^d$ ,  $t \in [T]$ . The interaction with the learner proceeds in  $T$  rounds as follows:

- ① In the  $t$ -th round, the  $t$ -th batch of coordinates  $x_t^{(1)}, \dots, x_t^{(n)} \in \mathbb{R}^d$  are revealed.
- ② After the  $t$ -th round, the algorithm is required to output  $\hat{\mu}_t \in \mathbb{R}^d$  at an estimate of  $\mu_t^*$ - the  $t$ -th batch of coordinates of  $\mu^*$ .

# Settings



At the end of this process, we say that the algorithm estimates the mean of  $\mathcal{X}$  under  $\varepsilon$ -corruption in the  $T$ -round online setting with error  $\varepsilon' > 0$ , failure probability  $\tau \in (0, 1)$  and sample complexity  $n$  if

$$\Pr[||\hat{\mu} - \mu^*||_2 - \varepsilon' \leq 0] = \Pr[\sqrt{\sum_{t=1}^T ||\hat{\mu}_t - \mu_t^*||_2^2} - \varepsilon' \leq 0] \geq 1 - \tau.$$

# Idea



This problem can be understood as calculating a weight for each sample. Suppose an algorithm generates a set of weights  $w_t^{(i)}$  for each sample  $i \in [n]$  at the end of round  $t$ . Then initialize the weights of the filtering algorithm in the  $(t+1)$ -th round to  $w_t^{(i)}$ . Unfortunately, the idea of simply applying the original weights and the current new input to update weights is theoretically not feasible. Consider the case where the unknown distribution  $X$  is the product of  $T$  isotropic Gaussians  $X_1, \dots, X_T$ . Then, the adversary can contaminate the sample set  $\mathcal{C}$  to make them look like *i.i.d.* samples from another isotropic Gaussian distribution  $X'_t$  at each round  $t \in [T]$ , such that  $\|E[X'_t] - E[X_t]\|_2 = c \cdot \varepsilon$  for some constant  $c$ . Then, the filtering algorithm should not downweight any sample, since the coordinates revealed in each round of contaminated samples are statistically indistinguishable from the *i.i.d.* samples from  $X'_t$ . Therefore, the algorithm's error still increases with  $\sqrt{T}$ , but the weights remain constant throughout.

# Algorithm




---

**Algorithm 1** Online Filter
 

---

```

1: Input: The number of samples  $n$ , Byzantine fraction  $\epsilon$ , round number  $T$ , sample coordinates  $x_t^{(1:n)}$ 
   revealed at the  $t^{\text{th}}$  round for  $t = 1, 2, \dots, T$ , filter threshold  $\lambda$ , and initialized weight  $w_0^{(i)} = 1/n$ , for
    $i = 1, 2, \dots, n$ .
2: for  $t = 1, 2, \dots, T$  do
3:   Initialize  $w_t \leftarrow w_{t-1}$  and update  $\tilde{x}_t^{(i)} = (x_1^{(i)}, \dots, x_t^{(i)})$ .
4:   Compute  $\Sigma \leftarrow \text{WCov}(w_t, \tilde{x}_t^{(1:n)})$ .
5:   while  $\|\Sigma\|_2 > 1 + \lambda$  do
6:     Compute the top eigenvector  $v$  of  $\Sigma$ .
7:     Compute empirical weighted mean  $\mu(w_t) \leftarrow \sum_{i=1}^n w_t^{(i)} \tilde{x}_t^{(i)}$ .
8:     for  $i = 1, 2, \dots, n$  do
9:       Compute  $\rho^{(i)} \leftarrow (v, \tilde{x}^{(i)} - \mu(w_t))^2$ .
10:    end for
11:    Sort  $\rho^{(1:n)}$  into a decreasing order denoted as  $\rho^{\pi(1)} \geq \rho^{\pi(2)} \geq \dots \geq \rho^{\pi(n)}$ , and let  $\beta$  be the smallest
    number such that  $\sum_{i=1}^{\beta} \rho^{\pi(i)} > 2\epsilon$ .
12:    Apply WFilter to update weights for  $\{\pi(1), \dots, \pi(\beta)\}$  as  $\{w_t^{\pi(1)}, \dots, w_t^{\pi(\beta)}\} \leftarrow$ 
    WFilter( $\rho^{\pi(1:\beta)}, w_t^{\pi(1:\beta)}$ ); while the remaining weights keep the same, i.e.,  $w_t^{\pi(i)} = w_t^{\pi(i)}$  for
     $i > \beta$ .
13:    Update  $\Sigma \leftarrow \text{WCov}(w_t, \tilde{x}_t^{(1:n)})$ 
14:  end while
15:  Output:  $\mu_t = \sum_{i=1}^n w_t^{(i)} \tilde{x}_t^{(i)}$ .
16: end for
  
```

---

**Subroutine 1:** Weighted Filter (WFilter)
 

---

```

1: Input: scores  $\rho^{\pi(1:\beta)}$ , weights  $w^{\pi(1:\beta)}$ .
2: for  $i = 1, 2, \dots, \beta$  do
3:    $w^{\pi(i)} \leftarrow (1 - \frac{\rho^{\pi(i)}}{\max_j \rho^{\pi(j)}}) w^{\pi(i)}$ .
4: end for
5: Output:  $w^{\pi(1:\beta)}$ .
  
```

---

**Subroutine2:** Weighted Covariance (WCov)
 

---

```

1: Input: weight  $w = w^{(1:n)}$  and samples  $\tilde{x}^{(1:n)}$ .
2: Compute  $\mu(w) \leftarrow \sum_{i=1}^n \frac{w^{(i)}}{\|w\|_1} \tilde{x}^{(i)}$ .
3: Compute weighted covariance estimation  $\Sigma \leftarrow \sum_{i=1}^n \frac{w^{(i)}}{\|w\|_1} (\tilde{x}^{(i)} - \mu(w))(\tilde{x}^{(i)} - \mu(w))^T$ .
4: Output:  $\Sigma$ .
  
```

---

# Result



## Definition 1.1.2 (( $\varepsilon, \delta$ )-stability)

For  $\varepsilon \in (0, \frac{1}{2})$  and  $\delta \geq \varepsilon$ , a finite set  $\mathcal{S} \subset \mathbb{R}^d$  is  $(\varepsilon, \delta)$ -stable with respect to a vector  $\mu \in \mathbb{R}^d$  if for every unit vector  $v \in \mathbb{R}^d$  and every subset  $\mathcal{S}' \subseteq \mathcal{S}$ , where  $|\mathcal{S}'| \geq (1-\varepsilon)|\mathcal{S}|$ , the following conditions are satisfied:

- 1  $\left| \frac{1}{|\mathcal{S}'|} \sum_{x \in \mathcal{S}'} v^T \cdot (x - \mu) \right| \leq \delta.$
- 2  $\left| \frac{1}{|\mathcal{S}'|} \sum_{x \in \mathcal{S}'} (v^T \cdot (x - \mu))^2 - 1 \right| \leq \frac{\delta^2}{\varepsilon}$

Suppose that  $\mathcal{C}$  is  $(\varepsilon, \delta)$ -stable with respect to  $\mu^*$  and  $\mathcal{X}$  is an  $\varepsilon$ -corrupted version of  $\mathcal{C}$ . Then, for  $\varepsilon$  at most a sufficiently small positive constant, there exists some constant  $\kappa$  such that their algorithm, when  $\lambda = \frac{\kappa \delta^2}{\varepsilon}$ , outputs a sequence of estimates satisfying  $\|\hat{\mu} - \mu^*\|_2 = O(\delta \log T).$



# Result



If the distribution is more structured, there is an algorithm that achieves the error with no dependence on  $T$  whatsoever.



Other topics



**Online Algorithms**

Online Robust Mean Estimation

Dynamic algorithms for k-center on graphs

Bin Packing under Random-Order: Breaking the barrier of  $\frac{3}{2}$

# Dynamic algorithms for k-center on graphs



# Settings



Given a metric space with  $n$  points and a positive integer  $k \leq n$ , the goal of the  $k$ -center problem is to select  $k$  points, referred to as centers, such that the maximum distance of any point in the metric space to its closest center is minimized.

The formal distribution is given below:

Given a weighted undirected graph  $G = (V, E, w)$  and an integer  $k \geq 1$ , the goal is to output a subset of vertices  $S \subseteq V$  of size at most  $k$ , such that the value  $\max_{v \in V} d_G(v, S)$  is minimized.

In the online setting, edges can be inserted into  $G$  or deleted from  $G$ . And the goal is to minimize the time complexity in each update and achieve better performance.

# Results



Given a weighted undirected graph  $G = (V, E, w)$  subject to edge updates, an integer parameter  $k \geq 1$ , and a positive constant parameter  $\varepsilon \leq \frac{1}{2}$ , there are two fully dynamic algorithms for the  $k$ -center problem on graphs, that maintain a  $(2 + \varepsilon)$ -approximation with the following guarantees (based on the current value of the matrix multiplication exponent):

- ① Deterministic algorithm with  $O(kn^{1.529}\varepsilon^{-2})$  worst-case update time, if  $G$  has uniform weights;
- ② Randomized algorithm, against an adaptive adversary, with  $O(kn^{1.823}\varepsilon^{-2})$  worst-case update time, if  $G$  has general weights.

Both algorithms have preprocessing time  $O(n^{2.373}\varepsilon^{-2} \log \varepsilon^{-1})$ .



## Results

Two special online setting are also concerned.

- 1 For the incremental setting, in which only edge insertions are allowed, given a weighted undirected graph  $G = (V, E, w)$  subject to edge insertions, an integer parameter  $k \geq 1$ , and a positive constant parameter  $\varepsilon < 1$ , there is a randomized incremental  $(4 + \varepsilon)$ -approximation algorithm for the  $k$ -center problem on graphs. The algorithm is correct *w.h.p.* and has  $kn^{o(1)}$  amortized update time over a sequence of  $\Theta(m)$  updates.
- 2 For the decremental setting, in which only edge deletions are allowed, given a weighted undirected graph  $G = (V, E, w)$  subject to edge deletions, an integer parameter  $k \geq 1$ , and a positive constant parameter  $\varepsilon < 1$ , there is a deterministic decremental  $(2 + \varepsilon)$ -approximation algorithm for the  $k$ -center problem on graphs, with  $kn^{o(1)}$  amortized update time over a sequence of  $\Theta(m)$  updates.



Other topics



**Online Algorithms**

Online Robust Mean Estimation

Dynamic algorithms for k-center on graphs

Bin Packing under Random-Order: Breaking the barrier of  $\frac{3}{2}$

# Bin Packing under Random-Order: Breaking the barrier





# Settings



Given a list  $I := (x_1, \dots, x_n)$ , where  $\forall i \in [n], x_i \in (0, 1]$ , the goal is to partition them into minimum number of unit size bins. In the online setting, item sizes are revealed one by one: in round  $i$ , the item  $x_i$  arrives and needs to be irrevocably assigned to a bin before the next items are revealed.

For a deterministic algorithm  $\mathcal{A}$ , the performance can be measured as the competitive ratio  $R_{\mathcal{A}}^{\infty} = \limsup_{m \rightarrow \infty} \left( \sup_{I: OPT(I)=m} \left( \frac{\mathcal{A}(I)}{OPT(I)} \right) \right)$ , where  $\mathcal{A}(I)$  denotes the number of bins used by  $\mathcal{A}$  to pack an input sequence  $I$ .

# Settings



For the online setting with the random-order model, the input set of items is chosen from adversary and the arrival order of the items is decided according to a permutation chosen uniformly at random from  $\mathcal{S}_n$ , the set of permutations of  $n$  elements. This reshuffling of the input items often weakens the adversary and provides better performance guarantees. In this model, the performance of an online algorithm  $\mathcal{A}$  can be measured using the following quantity, called random-order ratio ( $RR$ ):

$$RR_{\mathcal{A}}^{\infty} = \lim_{m \rightarrow \infty} \sup \left( \sup_{I: OPT(I)=m} \left( \frac{\mathbf{E}_{\sigma}[\mathcal{A}(I_{\sigma})]}{OPT(I)} \right) \right)$$

# History



Best-Fit algorithm is one of the most widely-used algorithms for bin packing. The idea is natural, simple and behaves well in practice. Best-Fit packs each item into the fullest bin where it fits, and open a new bin only if the item fits into none of the present opening bins.

It's proved that the  $RR$  of Best-Fit is upper-bounded by  $\frac{3}{2}$  and lower-bounded by 1.08. And in that paper, the author conjectured that the true ratio should lie somewhere close to 1.15.

# Results



- 1 Shows that Best-Fit achieves a  $RR$  of at most  $1.5 - \varepsilon$  for a small constant  $\varepsilon > 0$ .
- 2 Improve the lower bound of random order ratio to 1.144.

# Technique



Consider the last time  $t_\sigma$  of putting an item of size at most  $\frac{1}{3}$  in a bin of load at most  $\frac{1}{2}$ . Previous result shows that all bins except at most one that opened by Best-Fit to pack the first  $t_\sigma$  items are filled up to the level of at least  $\frac{2}{3}$ . And to pack items arriving after  $t_\sigma$ , Best-Fit is within a  $\frac{3}{2}$  factor of OPT. Combining these observations, an upper bound of  $\frac{3}{2}$  was achieved.

To break the barrier, they did further analysis that either the factor of the period before  $t_\sigma$  can be improved to  $\frac{3}{2} - 2\varepsilon$ , or the factor after  $t_\sigma$  can be improved to  $\frac{3}{2} - 2\varepsilon$ . They did a case analysis based on  $t_\sigma$  and show that either a large fraction of the bin packed by Best-Fit is rather full (the load is at least  $\frac{3}{4}$ ) or Best-Fit performs relatively well compared to OPT.

# 谢谢



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY

